

Chapter 12 :: Concurrency

Programming Language Pragmatics

Michael L. Scott

Background and Motivation

- A PROCESS or THREAD is a potentially-active execution context
- Classic von Neumann (stored program) model of computing has single thread of control
- Parallel programs have more than one
- A process can be thought of as an abstraction of a physical PROCESSOR

Background and Motivation

- Processes/Threads can come from
 - multiple CPUs
 - kernel-level multiplexing of single physical machine
 - language or library level multiplexing of kernel-level abstraction
- They can run
 - in true parallel
 - unpredictably interleaved
 - run-until-block
- Most work focuses on the first two cases, which are equally difficult to deal with

Background and Motivation

- Two main classes of programming notation
 - synchronized access to shared memory
 - message passing between processes that don't share memory
- Both approaches can be implemented on hardware designed for the other, though shared memory on message-passing hardware tends to be slow

Background and Motivation

- Race conditions
 - A race condition occurs when actions in two processes are not synchronized and program behavior depends on the order in which the actions happen
 - Race conditions are not all bad; sometimes any of the possible program outcomes are ok (e.g. workers taking things off a task queue)

Background and Motivation

- Race conditions (we want to avoid race conditions):
 - Suppose processors A and B share memory, and both try to increment variable X at more or less the same time
 - Very few processors support arithmetic operations on memory, so each processor executes
 - LOAD X
 - INC
 - STORE X
 - Suppose X is initialized to 0. If both processors execute these instructions simultaneously, what are the possible outcomes?
 - could go up by one or by two

Background and Motivation

- Synchronization
 - SYNCHRONIZATION is the act of ensuring that events in different processes happen in a desired order
 - Synchronization can be used to eliminate race conditions
 - In our example we need to synchronize the increment operations to enforce MUTUAL EXCLUSION on access to X
 - Most synchronization can be regarded as either
 - Mutual exclusion (making sure that only one process is executing a CRITICAL SECTION [touching a variable, for example] at a time),
or as
 - CONDITION SYNCHRONIZATION, which means making sure that a given process does not proceed until some condition holds (e.g. that a variable contains a given value)

Background and Motivation

- One might be tempted to think of mutual exclusion as a form of condition synchronization (the condition being that nobody else is in the critical section), but it isn't
 - The distinction is basically *existential* v. *universal* quantification
 - Mutual exclusion requires multi-process consensus
- We do NOT in general want to over-synchronize
 - That eliminates parallelism, which we generally want to encourage for performance
- Basically, we want to eliminate "bad" race conditions, i.e., the ones that cause the program to give incorrect results

Background and Motivation

- Historical development of shared memory ideas
 - To implement synchronization you have to have something that is ATOMIC
 - that means it happens all at once, as an indivisible action
 - In most machines, reads and writes of individual memory locations are atomic (note that this is not trivial; memory and/or busses must be designed to arbitrate and serialize concurrent accesses)
 - In early machines, reads and writes of individual memory locations were *all* that was atomic
 - To simplify the implementation of mutual exclusion, hardware designers began in the late 60's to build so-called read-modify-write, or fetch-and-phi, instructions into their machines

Concurrent Programming Fundamentals

- SCHEDULERS give us the ability to "put a thread/process to sleep" and run something else on its process/processor
 - start with coroutines
 - make uniprocessor run-until-block threads
 - add preemption
 - add multiple processors

Concurrent Programming Fundamentals




Figure 12.5 Lifetime of concurrent threads. With co-begin, parallel loops, or launch-at-elaboration (a), threads are always properly nested. With fork/join (b), more general patterns are possible.

Concurrent Programming Fundamentals

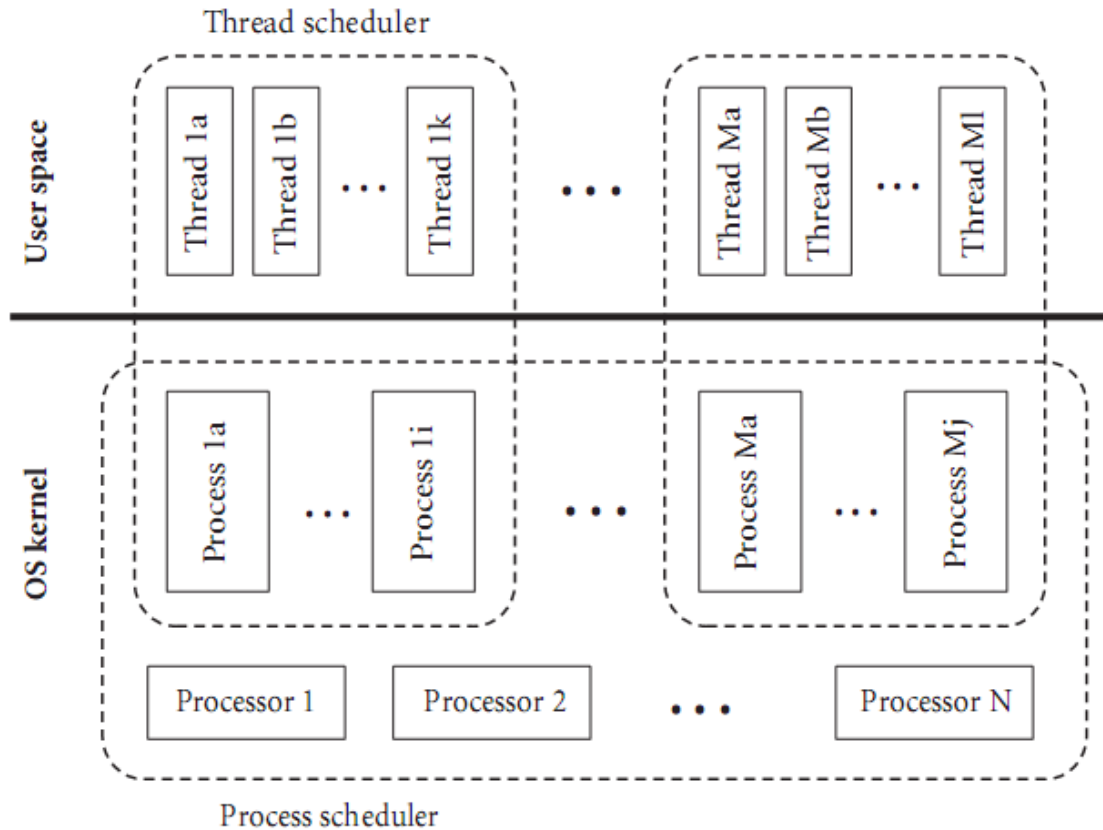


Figure 12.6 Two-level implementation of threads. A thread scheduler, implemented in a library or language run-time package, multiplexes threads on top of one or more kernel-level processes, just as the process scheduler, implemented in the operating system kernel, multiplexes processes on top of one or more physical processors.

Concurrent Programming Fundamentals

- Coroutines
 - Multiple execution contexts, only one of which is active
 - Transfer (other):
 - save all callee-saves registers on stack, including ra and fp
 - `*current := sp`
 - `current := other`
 - `sp := *current`
 - pop all callee-saves registers (including ra, but NOT sp!)
 - return (into different coroutine!)
 - Other and current are pointers to CONTEXT BLOCKs
 - Contains sp; may contain other stuff as well (priority, I/O status, etc.)
 - No need to change PC; always changes at the same place
 - Create new coroutine in a state that looks like it's blocked in transfer.
(Or maybe let it execute and then "detach". That's basically early reply)

Concurrent Programming Fundamentals

- Run-until block threads on a single process
 - Need to get rid of explicit argument to transfer
 - Ready list data structure: threads that are runnable but not running

```
procedure reschedule:
    t : cb := dequeue(ready_list)
    transfer(t)
```
 - To do this safely, we need to save 'current' somewhere - two ways to do this:
 - Suppose we're just relinquishing the processor for the sake of fairness (as in MacOS or Windows 3.1):

```
procedure yield:
    enqueue (ready_list, current)
    reschedule
```
 - Now suppose we're implementing synchronization:

```
sleep_on(q)
    enqueue(q, current)
    reschedule
```
 - Some other thread/process will move us to the ready list when we can continue

Concurrent Programming Fundamentals

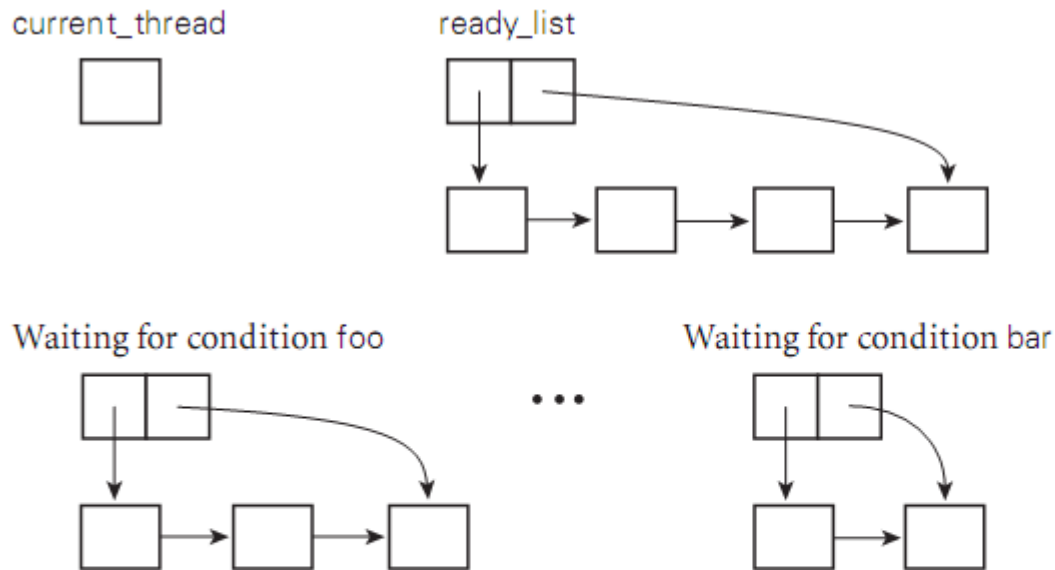


Figure 12.7 Data structures of a simple scheduler. A designated `current_thread` is running. Threads on the ready list are runnable. Other threads are blocked, waiting for various conditions to become true. If threads run on top of more than one OS-level process, each such process will have its own `current_thread` variable. If a thread makes a call into the operating system, its process may block in the kernel.

Concurrent Programming Fundamentals

- Preemption

- Use timer interrupts (in OS) or signals (in library package) to trigger involuntary yields

- Requires that we protect the scheduler data structures:

```
procedure yield:
    disable_signals
    enqueue(ready_list, current)
    Reschedule
    re-enable_signals
```

- Note that reschedule takes us to a different thread, possibly in code other than yield Invariant: EVERY CALL to reschedule must be made with signals disabled, and must re-enable them upon its return

```
    disable_signals
    if not <desired condition>
        sleep_on <condition queue>
    re-enable_signals
```


Concurrent Programming Fundamentals

- Multiprocessors

- Disabling signals doesn't suffice:

```
procedure yield:
    disable_signals
    acquire(scheduler_lock)           // spin lock
    enqueue(ready_list, current)
    reschedule
    release(scheduler_lock)
    re-enable_signals
```

```
disable_signals
acquire(scheduler_lock)           // spin lock
if not <desired condition>
    sleep_on <condition queue>
release(scheduler_lock)
re-enable_signals
```

Implementing Synchronization

- Condition synchronization with atomic reads and writes is easy
 - You just cast each condition in the form of "location X contains value Y" and you keep reading X in a loop until you see what you want
- Mutual exclusion is harder
 - Much early research was devoted to figuring out how to build it from simple atomic reads and writes
 - Dekker is generally credited with finding the first correct solution for two processes in the early 1960s
 - Dijkstra published a version that works for N processes in 1965
 - Peterson published a much simpler two-process solution in 1981

Implementing Synchronization

- Repeatedly reading a shared location until it reaches a certain value is known as SPINNING or BUSY-WAITING
- A busy-wait mutual exclusion mechanism is known as a SPIN LOCK
 - The problem with spin locks is that they waste processor cycles
 - Synchronization mechanisms are needed that interact with a thread/process scheduler to put a process to sleep and run something else instead of spinning
 - Note, however, that spin locks are still valuable for certain things, and are widely used
 - In particular, it is better to spin than to sleep when the expected spin time is less than the rescheduling overhead

Implementing Synchronization

- SEMAPHORES were the first proposed SCHEDULER-BASED synchronization mechanism, and remain widely used
- CONDITIONAL CRITICAL REGIONS and MONITORS came later
- Monitors have the highest-level semantics, but a few sticky semantic problem - they are also widely used
- Synchronization in Java is sort of a hybrid of monitors and CCRs (Java 3 will have true monitors.)
- Shared-memory synch in Ada 95 is yet another hybrid

Implementing Synchronization

- A semaphore is a special counter
- It has an initial value and two operations, P and V, for changing that value
- A semaphore keeps track of the difference between the number of P and V operations that have occurred
- A P operation is delayed (the process is de-scheduled) until $\#P - \#V \leq C$, the initial value of the semaphore

Implementing Synchronization

```
shared scheduler_lock : low_level_lock
shared ready_list : queue of thread
per-process private current_thread : thread

procedure reschedule
  -- assume that scheduler_lock is already held
  -- and that timer signals are disabled
  t : thread
  loop
    t := dequeue(ready_list)
    if t ≠ null
      exit
    -- else wait for a thread to become runnable
    release_lock(scheduler_lock)
    -- window allows another thread to access ready_list
    -- (no point in reenabling signals;
    -- we're already trying to switch to a different thread)
    acquire_lock(scheduler_lock)
  transfer(t)
  -- caller must release scheduler_lock
  -- and reenable timer signals after we return

procedure yield
  disable_signals
  acquire_lock(scheduler_lock)
  enqueue(ready_list, current_thread)
  reschedule
  release_lock(scheduler_lock)
  reenable_signals

procedure sleep_on(ref Q : queue of thread)
  -- assume that caller has already disabled timer signals
  -- and acquired scheduler_lock, and will reverse
  -- these actions when we return
  enqueue(Q, current_thread)
  reschedule
```

Note: a possible implementation is shown on the next slide

Figure 12.12 Pseudocode for part of a simple reentrant (parallelism-safe) scheduler. Every process has its own copy of `current_thread`. There is a single shared `scheduler_lock` and a single `ready_list`. If processes have dedicated processors, then the `low_level_lock` can be an ordinary spin lock; otherwise it can be a “spin-then-yield” lock (Figure 12.13). The loop inside `reschedule` busy-waits until the ready list is nonempty. The code for `sleep_on` cannot disable timer signals and acquire the scheduler lock itself, because the caller needs to test a condition and then block as a single atomic operation.

Implementing Synchronization

```
type semaphore = record
  N : integer -- usually initialized to something nonnegative
  Q : queue of threads

procedure P(ref S : semaphore)
  disable_signals
  acquire_lock(scheduler_lock)
  S.N -= 1
  if S.N < 0
    sleep_on(S.Q)
  release_lock(scheduler_lock)
  reenable_signals

procedure V(ref S : semaphore)
  disable_signals
  acquire_lock(scheduler_lock)
  S.N += 1
  if N ≤ 0
    -- at least one thread is waiting
    enqueue(ready_list, dequeue(S.Q))
  release_lock(scheduler_lock)
  reenable_signals
```

Figure 12.14 Semaphore operations, for use with the scheduler code of Figure 12.12.

Implementing Synchronization

```
shared buf : array [1..SIZE] of bdata
shared next_full, next_empty : integer := 1, 1
shared mutex : semaphore := 1
shared empty_slots, full_slots : semaphore := SIZE, 0

procedure insert(d : bdata)
    P(empty_slots)
    P(mutex)
    buf[next_empty] := d
    next_empty := next_empty mod SIZE + 1
    V(mutex)
    V(full_slots)

function remove : bdata
    P(full_slots)
    P(mutex)
    d : bdata := buf[next_full]
    next_full := next_full mod SIZE + 1
    V(mutex)
    V(empty_slots)
    return d
```

Figure 12.15 Semaphore-based code for a bounded buffer. The `mutex` binary semaphore protects the data structure proper. The `full_slots` and `empty_slots` general semaphores ensure that no operation starts until it is safe to do so.

Implementing Synchronization

- It is generally assumed that semaphores are fair, in the sense that processes complete P operations in the same order they start them
- Problems with semaphores
 - They're pretty low-level.
 - When using them for mutual exclusion, for example (the most common usage), it's easy to forget a P or a V, especially when they don't occur in strictly matched pairs (because you do a V inside an if statement, for example, as in the use of the spin lock in the implementation of P)
 - Their use is scattered all over the place.
 - If you want to change how processes synchronize access to a data structure, you have to find all the places in the code where they touch that structure, which is difficult and error-prone

Language-Level Mechanisms

- Monitors were an attempt to address the two weaknesses of semaphores listed above
- They were suggested by Dijkstra, developed more thoroughly by Brinch Hansen, and formalized nicely by Hoare (a real cooperative effort!) in the early 1970s
- Several parallel programming languages have incorporated monitors as their fundamental synchronization mechanism
 - none, incorporates the precise semantics of Hoare's formalization

Language-Level Mechanisms

- A monitor is a shared object with operations, internal state, and a number of condition queues. Only one operation of a given monitor may be active at a given point in time
- A process that calls a busy monitor is delayed until the monitor is free
 - On behalf of its calling process, any operation may suspend itself by waiting on a condition
 - An operation may also signal a condition, in which case one of the waiting processes is resumed, usually the one that waited first

Language-Level Mechanisms

- The precise semantics of mutual exclusion in monitors are the subject of considerable dispute. Hoare's original proposal remains the clearest and most carefully described
 - It specifies two bookkeeping queues for each monitor: an entry queue, and an urgent queue
 - When a process executes a signal operation from within a monitor, it waits in the monitor's urgent queue and the first process on the appropriate condition queue obtains control of the monitor
 - When a process leaves a monitor it unblocks the first process on the urgent queue or, if the urgent queue is empty, it unblocks the first process on the entry queue instead

Language-Level Mechanisms

- Building a correct monitor requires that one think about the "monitor invariant". The monitor invariant is a predicate that captures the notion "the state of the monitor is consistent."
 - It needs to be true initially, and at monitor exit
 - It also needs to be true at every wait statement
 - In Hoare's formulation, needs to be true at every signal operation as well, since some other process may immediately run
- Hoare's definition of monitors in terms of semaphores makes clear that semaphores can do anything monitors can
- The inverse is also true; it is trivial to build a semaphores from monitors (Exercise)